

The dominant sequence transduction models are based on complex recurrent or convolutional neural networks. We propose the Transformer, a model architecture relying entirely on an attention mechanism. The Transformer attention mechanism and BERT language model have advanced natural language processing through semantic similarity embeddings and fine-tuning of pre-trained models.